

Data Analysis

Ariane Stark

4/28/2021

```
dat <- read_csv("case_and_demographics.csv") %>%
  subset(select = -c(1)) %>%
  subset(select = c(1,3:4,2,5:6,43,56:61 )) %>%
  mutate(Cum_Case_Median = cut_number(cum_case,2, labels= c(0,1))) %>%
  na.omit() #removes Rio Arriba County NM which has NA in the pov. dem.
```

Cross Validation of Best Subset, Forward and Backward Stepwise

```
predict.regsubsets <- function (object ,newdata ,id ,...){
  form<-as.formula(object$call [[2]])
  mat<-model.matrix(form,newdata)
  coefi<-coef(object ,id=id)
  xvars<-names(coefi)
  mat[,xvars]%*%coefi
}
```

```
set.seed(100)

k<-10

folds<-sample(1:k,nrow(dat[,c(4:13)]),replace=TRUE)
cv.errors.best_subset <- matrix(NA,k,9, dimnames=list(NULL, paste(1:9)))
cv.errors.forward <- matrix(NA,k,9, dimnames=list(NULL, paste(1:9)))
cv.errors.backward <- matrix(NA,k,9, dimnames=list(NULL, paste(1:9)))

for(j in 1:k){
  # best subset select
  best_subset.fit <- regsubsets(cum_case~.,data = dat[folds!=j,c(4:13)],
                                method = "exhaustive", nvmax = 9)

  # forward stepwise select
  forward.fit <- regsubsets(cum_case~.,data = dat[folds!=j,c(4:13)],
                            method = "forward", nvmax = 9)

  # backward stepwise select
  backward.fit <- regsubsets(cum_case~.,data = dat[folds!=j,c(4:13)],
                            method = "backward", nvmax = 9)

  for(i in 1:9){
```

```

# best subset error
pred.best <-
  predict.regsbsets(best_subset.fit, dat[folds==j,c(4:13)], id=i)
cv.errors.best_subset[j,i] <-
  mean((dat$cum_case[folds==j]-pred.best)^2)

# forward stepwise error
pred.forward <-
  predict.regsbsets(forward.fit, dat[folds==j,c(4:13)], id=i)
cv.errors.forward[j,i] <-
  mean((dat$cum_case[folds==j]-pred.forward)^2)

# backward stepwise error
pred.backward <-
  predict.regsbsets(backward.fit, dat[folds==j,c(4:13)], id=i)
cv.errors.backward[j,i] <-
  mean((dat$cum_case[folds==j]-pred.backward)^2)

} }

```

```

#best subset CV model selection
mean.cv.errors.best <- apply(cv.errors.best_subset,2,mean)
mean.cv.errors.best

```

```


##          1          2          3          4          5          6          7          8
## 68750976 45843477 58650271 57273927 68389310 63810273 67377272 67568839
##          9
## 67254330

```

```

plot(mean.cv.errors.best,type="b",
     xlab="Number of Variables",
     ylab="Best Subset Mean C.V. Error")
points(which.min(mean.cv.errors.best),
       mean.cv.errors.best[which.min(mean.cv.errors.best)],
       col="red", cex=1.5,pch=20)

```




```
#forward stepwise CV model selection
```

```
mean.cv.errors.forward <- apply(cv.errors.forward,2,mean)
mean.cv.errors.forward
```

```
##      1      2      3      4      5      6      7      8
## 68750976 45843477 56429254 57273927 59513523 67205055 70711271 67190749
##      9
## 67254330
```

```
plot(mean.cv.errors.forward,type="b",
     xlab="Number of Variables",
     ylab="Forward Stepwise Mean C.V. Error")
points(which.min(mean.cv.errors.forward),
       mean.cv.errors.forward[which.min(mean.cv.errors.forward)],
       col="red", cex=1.5,pch=20)
```




```
#backward stepwise CV model selection
```

```
mean.cv.errors.backward <- apply(cv.errors.backward,2,mean)
```

```
mean.cv.errors.backward
```

```
##      1      2      3      4      5      6      7      8
## 129210484 128813623 93146095 65815873 68389310 63810273 67377272 67568839
##      9
## 67254330
```

```
plot(mean.cv.errors.backward,type="b",
     xlab="Number of Variables",
     ylab="Backward Stepwise Mean C.V. Error")
points(which.min(mean.cv.errors.backward),
       mean.cv.errors.backward[which.min(mean.cv.errors.backward)],
       col="red", cex=1.5,pch=20)
```



Best Subset Overall

```
best_subset.fit <- regsubsets(cum_case~., data = dat[,c(4:13)],
                             method = "exhaustive", nvmax = 9)
best_subset.summary <- summary(best_subset.fit)
best_subset.summary.frame <- data_frame("Parameters" = seq(1:9),
                                         "R^2" = best_subset.summary$rsq,
                                         "AdjR^2" = best_subset.summary$adjr2,
                                         "CP" = best_subset.summary$cp,
                                         "BIC" = best_subset.summary$bic)
```

```
## Warning: 'data_frame()' was deprecated in tibble 1.1.0.
## Please use 'tibble()' instead.
```

```
best_subset.summary.frame
```

```
## # A tibble: 9 x 5
##   Parameters 'R^2' 'AdjR^2'    CP    BIC
##   <int> <dbl>    <dbl> <dbl> <dbl>
## 1     1 0.945    0.945 2488. -9330.
## 2     2 0.964    0.964  546. -10664.
## 3     3 0.966    0.966  306. -10870.
```

```
## 4      4 0.967    0.967 178.   -10983.
## 5      5 0.968    0.968  89.5  -11063.
## 6      6 0.969    0.969   5.65 -11139.
## 7      7 0.969    0.969   6.14 -11133.
## 8      8 0.969    0.969   8.07 -11125.
## 9      9 0.969    0.969  10.0  -11117.
```

Based on Cross Validation we pick the model with 2 parameters

```
coef(best_subset.fit,2)
```

```
##      (Intercept)      Total_Pop Total_Hispanic
## 315.11931521      0.07300349      0.08304726
```

Forward Stepwise Overall

```
forward.fit <- regsubsets(cum_case~.,data = dat[,c(4:13)],
                          method = "forward", nvmax = 9)
forward.summary <- summary(forward.fit)
forward.summary.frame <- data_frame("Parameters" = seq(1:9),
                                    "R^2"=forward.summary$rsq,
                                    "AdjR^2"=forward.summary$adjr2,
                                    "CP"=forward.summary$cp,
                                    "BIC"=forward.summary$bic)

forward.summary.frame
```

```
## # A tibble: 9 x 5
##   Parameters 'R^2' 'AdjR^2'    CP    BIC
##   <int> <dbl>   <dbl>   <dbl> <dbl>
## 1     1 0.945   0.945 2488.  -9330.
## 2     2 0.964   0.964  546. -10664.
## 3     3 0.966   0.966  306. -10870.
## 4     4 0.967   0.967  178. -10983.
## 5     5 0.967   0.967  176. -10979.
## 6     6 0.969   0.969   5.65 -11139.
## 7     7 0.969   0.969   6.14 -11133.
## 8     8 0.969   0.969   8.07 -11125.
## 9     9 0.969   0.969  10.0 -11117.
```

Based on Cross Validation we pick the model with 2 parameters

```
coef(forward.fit,2)
```

```
##      (Intercept)      Total_Pop Total_Hispanic
## 315.11931521      0.07300349      0.08304726
```

Backward Stepwise Overall

```
backward.fit <- regsubsets(cum_case~.,data = dat[,c(4:13)],
                          method = "backward", nvmax = 9)
backward.summary <- summary(backward.fit)
backward.summary.frame <- data_frame("Parameters" = seq(1:9),
                                     "R^2"=backward.summary$rsq,
                                     "AdjR^2"=backward.summary$adjr2,
                                     "CP"=backward.summary$cp,
                                     "BIC"=backward.summary$bic)

backward.summary.frame
```

```
## # A tibble: 9 x 5
##   Parameters 'R^2' 'AdjR^2'    CP    BIC
##   <int> <dbl>    <dbl>    <dbl>    <dbl>
## 1     1 0.903    0.903 6923.    -7478.
## 2     2 0.956    0.956 1375.   -10022.
## 3     3 0.965    0.965  452.   -10739.
## 4     4 0.967    0.967  271.   -10896.
## 5     5 0.968    0.968   89.5  -11063.
## 6     6 0.969    0.969    5.65 -11139.
## 7     7 0.969    0.969    6.14 -11133.
## 8     8 0.969    0.969    8.07 -11125.
## 9     9 0.969    0.969   10.0  -11117.
```

Based on Cross Validation we pick the model with 6 parameters

```
coef(best_subset.fit,6)
```

```
##              (Intercept)              Total_Pop
##          -338.60412493          -0.38868473
##              Total_White              Total_Black
##          0.58210948          0.55858828
##              Total_Hispanic              Total_Other
##          0.05514001          0.51284419
## Total_Households_Above_Poverty
##          -0.29137162
```

Ridge and Lasso

```
set.seed(100)
k <- 10
cv.error.lasso <- rep(NA,k)
cv.error.ridge <- rep(NA,k)
cv.lam.lasso <- rep(NA,k)
cv.lam.ridge <- rep(NA,k)
set.seed(1)
folds<-sample(1:k,nrow(dat[,c(4:13)]),replace=TRUE)
for (i in 1:k){
  train <- dat[folds!=i,c(4:13)] #Set the training set
  test <- dat[folds==i,c(4:13)] #Set testing set
```

```

x.train <- model.matrix(cum_case~., data=train)[,-1]
y.train <- train$cum_case
x.test <- model.matrix(cum_case~., data=test)[,-1]
y.test <- test$cum_case

#Lasso
lasso.mod <- glmnet(x.train,y.train,alpha=1)
cv.out <- cv.glmnet(x.train,y.train,alpha=1)
bestlam <- cv.out$lambda.min
cv.lam.lasso[i] <- bestlam
lasso.pred <- predict(lasso.mod,s=bestlam ,newx=x.test)
cv.error.lasso[i] <- mean((lasso.pred-y.test)^2)

#Ridge
ridge.mod <- glmnet(x.train,y.train,alpha=0)
cv.out <- cv.glmnet(x.train,y.train,alpha=0)
bestlam <- cv.out$lambda.min
cv.lam.ridge[i] <- bestlam
ridge.pred <- predict(ridge.mod,s=bestlam ,newx=x.test)
cv.error.ridge[i] <- mean((ridge.pred-y.test)^2)
}

```

```

# Lasso CV Errors and Mean
cv.error.lasso[which.min(cv.error.lasso)]

```

```
## [1] 21711998
```

```
mean(cv.error.lasso)
```

```
## [1] 64144413
```

```

# Ridge CV Errors and Mean
cv.error.ridge[which.min(cv.error.ridge)]

```

```
## [1] 24901393
```

```
mean(cv.error.ridge)
```

```
## [1] 58650401
```

Ridge Regression performs better when the response is a function of many predictors of roughly equal size

Ridge With Best CV Error Lambda


```
x<-model.matrix(cum_case~., data=dat[,c(4:13)])[, -1]
y<-dat$cum_case

ridge.fit<-glmnet(x,y,alpha=0)
ridge.coef<-predict(ridge.fit,type="coefficients",
                    s=cv.lam.ridge[which.min(cv.error.ridge)])[1:9,]
ridge.coef[ridge.coef!=0]
```

```
##                (Intercept)                POP_DENSITY
##                2659.26735391                -1.46826014
##                Total_Pop                Median_Age
##                0.01675252                -60.01296350
##                Total_White                Total_Black
##                0.03389108                0.04046944
##                Total_Hispanic                Total_Other
##                0.06843143                0.03043294
## Total_Households_Below_Poverty
##                0.27874317
```

Lasso With Best CV Error Lambda

```
lasso.fit<-glmnet(x,y,alpha=1)
lasso.coef<-predict(lasso.fit,type="coefficients",
                    s=cv.lam.lasso[which.min(cv.error.lasso)])[1:9,]
lasso.coef[lasso.coef!=0]
```

```
##                (Intercept)                Total_Pop
##                286.95400883                0.04444873
##                Total_White                Total_Black
##                0.02943025                0.01685243
##                Total_Hispanic Total_Households_Below_Poverty
##                0.08547591                0.09854572
```

Regression Trees

```
tree_dat <- dat %>%
  subset(select = c(4:13))

tree <- rpart(cum_case ~.,
              data=tree_dat,control=rpart.control(cp=.0001))

#printcp(tree)

best <- tree$scptable[which.min(tree$scptable[, "xerror"]), "CP"]
pruned_tree <- prune(tree, cp=best)
prp(pruned_tree,
    faclen=0, #use full names for factor labels
```

```

extra=1, #display number of obs. for each terminal node
roundint=F, #don't round to integers in output
digits=5) #display 5 decimal places in output

```

